

Learning to Predict and Improve Build Successes in Package Ecosystems



Harshitha Menon, Todd Gamblin
Lawrence Livermore National Laboratory

BUILD-SI



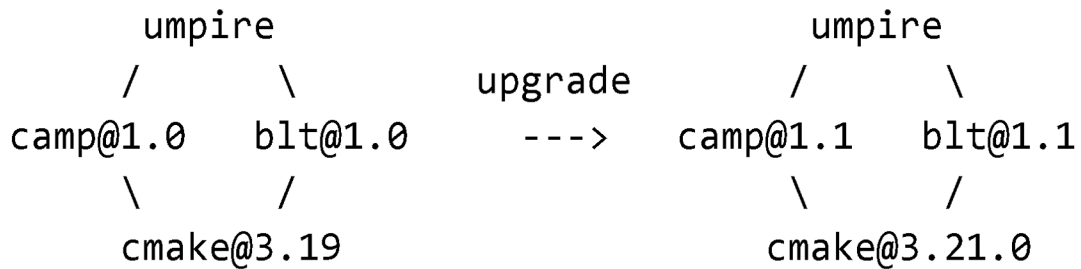
We build codes from hundreds of small, complex pieces

- Component-based software development dates back to the 60's
- M.D. McIlroy, Mass Produced Software Components. NATO SE Conf., 1968
- Pros are well known:
 - Teams can and must reuse each others' work
 - Teams write less code, meet deliverables faster
- Cons:
 - Teams must ensure that components work together
 - Integration burden increases with each additional library
 - Integration must be repeated with each update to components

Developers integrate large software stacks manually

- Simply finding a compatible set of versions for packages is NP-hard (combinatorial)
 - Changing one version may affect all others, those changes may trigger others, and so on
- In HPC, there are many more parameters to adjust
 - Version, compiler, ABI, build options, microarchitecture, GPU capability, etc.
 - Multiple codes in the same environment (workflows), performance
- We solve this problem repeatedly by trial and error
 - Incompatibilities are not known in advance; developers discover them

Transitive dependency requirements can cause cascading errors



- blt@1.0 requires cmake ≥ 3.18 , but is incompatible with cmake@3.21.0 due to an unknown bug
- camp@1.0 depends on cmake@3.19 or higher, but camp@1.1 depends on cmake@3.21 or higher
- The umpire developers want to use camp@1.1 for its new features
- Upgrading camp to 1.1 pushes cmake to the latest 3.21.0 will cause the build to fail
- We need to use blt@1.1 to make this work.

Team would have to build several versions of cmake and blt to find a working configuration

BuildCheck

Predict the build outcome of various package configurations with high accuracy.

Graph Neural Networks (GNN) for Build Prediction


Why GNN?

- GNNs are ideally suited for handling complex dependency relationships in software ecosystems.

Problem Definition

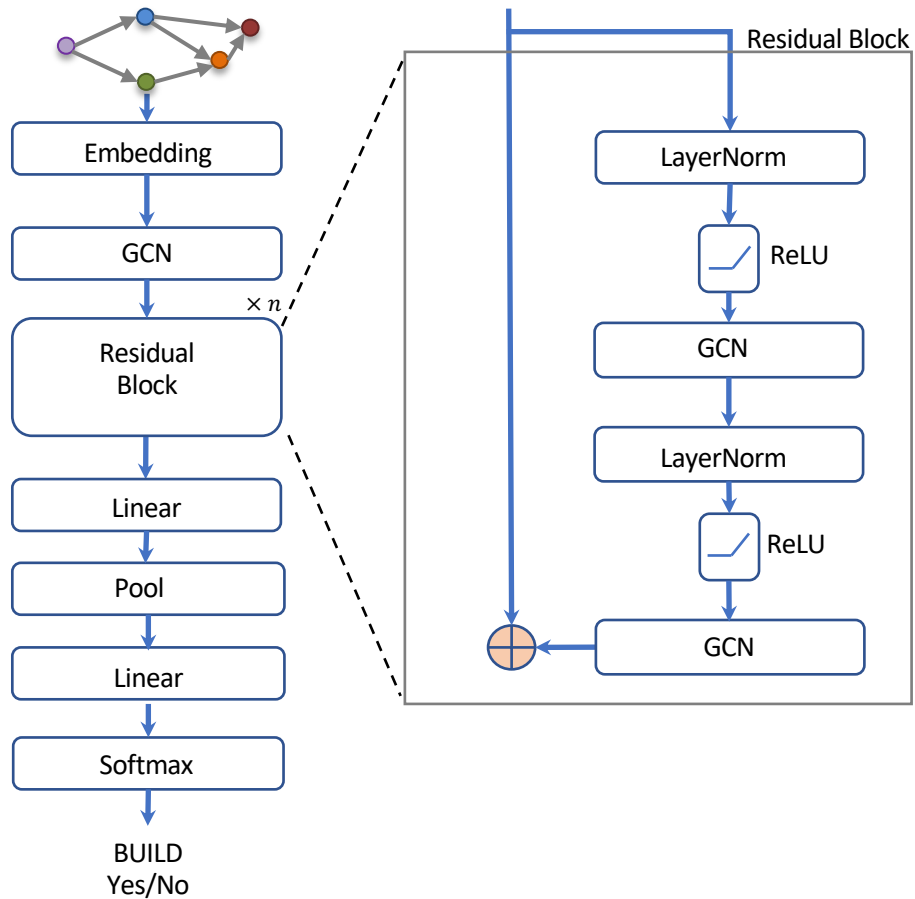
- The package dependency graph is a directed acyclic graph (DAG)
- Graph is represented as $G = (V, E)$, where V is the set of nodes representing the packages and E is the set of edges capturing the dependencies.
- We cast the build success prediction problem as a supervised learning problem
- **Goal:** learn a model to predict the build outcome

Overview of build outcome prediction using GNN



- Each configuration is represented as a graph
- Node features incorporate information about packages (which package and version)
- Layers GCN
- Final layer does a global pooling to predict whether this configuration builds or not.


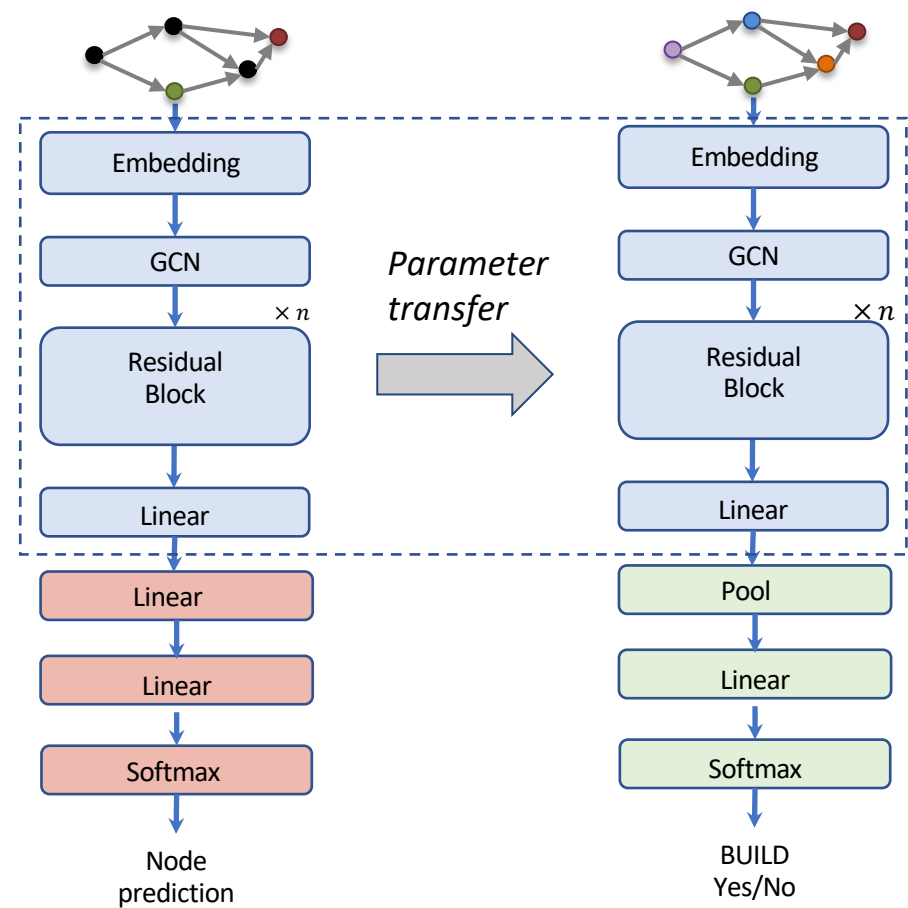
BuildCheck GNN Architecture



Main Components


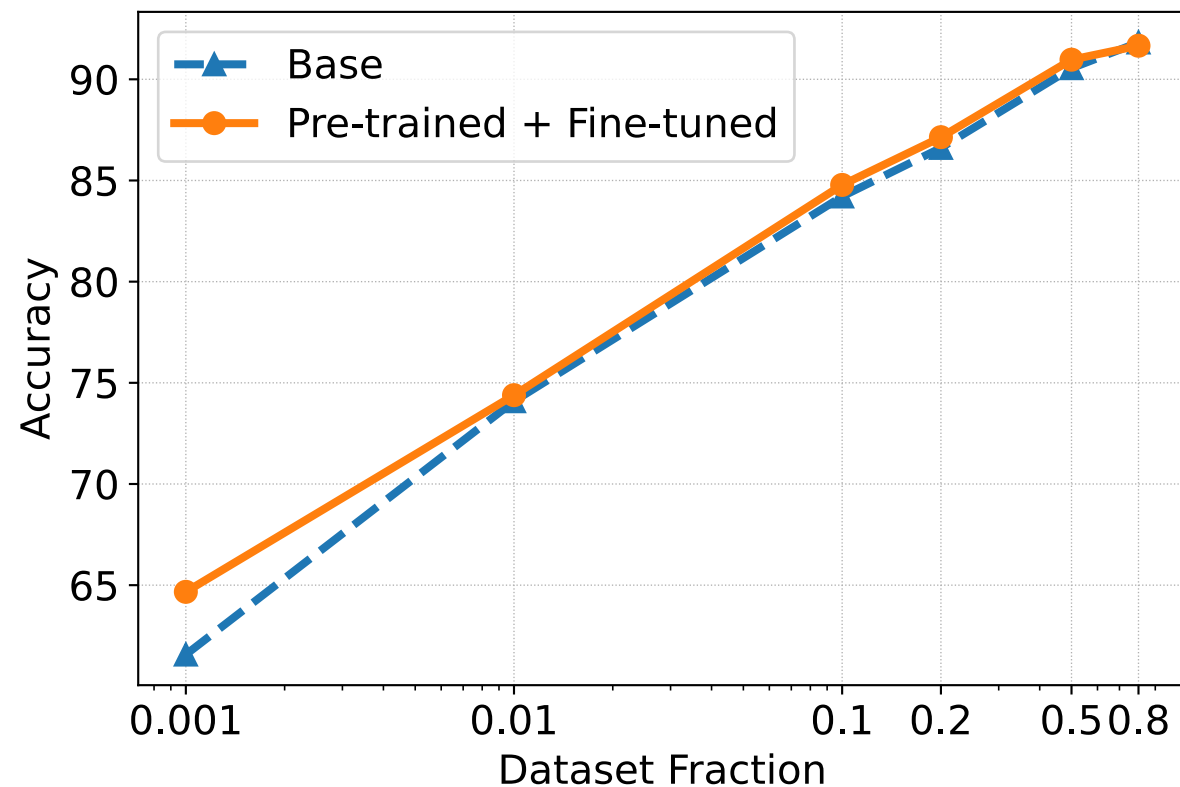
- Multiple Graph Convolutional layers.
- Embedding layer: maps package information into a continuous vector space
- Residual block: aids in training deeper networks.
- Pool layer: computes the average of all node features and creates a representation of the entire graph

Self-Supervised Pretraining Task for Learning Node Embeddings for Downstream tasks



Pre-trained package dependency model can be used for build prediction with fine-tuning

Evaluation



Our model achieves an accuracy of 91% on E4S build dataset

Conclusion and Ongoing Work

- Demonstrated how to combine the capabilities of Graph Neural Networks and advanced package management technologies to offer practical solutions for managing package dependencies
- BuildCheck can eliminate very expensive trial-and-error exercise to find working builds
- Use the outputs of BuildCheck in Spack's concretizer solver
- Using GNN Explainer to explain the reason for build failure prediction
- Integrate the model with the Spack's Continuous Integration system



CASC

Center for Applied
Scientific Computing



Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

This work was prepared by LLNL under Contract DE-AC52-07NA27344 and was supported by the LLNL-LDRD Program under Project No. 21-SI-005.